AI-Based Fraud Detection System: Machine Learning Approaches for Detecting Fraudulent Transactions in Digital Payment Systems

Mohd Saif
Department of Computer Science & Engineering
Noida Institute of Engineering and Technology
(of Affiliation)
Dr. APJ Abdul Kalam University
Gr. Noida, India
[mdsaif979550@gmail.com]

Devbrat Prasad
Department of Computer Science & Engineering
Noida Institute of Engineering and Technology
(of Affiliation)
Dr. APJ Abdul Kalam University
Gr. Noida, India
[devbrat242@gmail.com]

Gaurav Sharma
Department of Computer Science & Engineering
Noida Institute of Engineering and Technology
(of Affiliation)
Dr. APJ Abdul Kalam University
Gr. Noida, India
[gsharma.dn9115@gmail.com]

Ishant Thakur
Department of Computer Science & Engineering
Noida Institute of Engineering and Technology
(of Affiliation)
Dr. APJ Abdul Kalam University
Gr. Noida, India
[ishantthakur958@gmail.com]

Supervisor:

Ms.Garima Jain

Deputy Head,CSE

Department of Computer Science and Engineering

Noida Institute of Engineering and Technology,Gr.Noida

[garimajain@niet.co.in]

***Abstract*—** The rise of digital payment platforms has significantly increased the risk of fraudulent transactions, creating a need for more adaptive and intelligent detection systems. This study presents the development of a machine learning-based fraud detection model that classifies transactions as legitimate or fraudulent using Logistic Regression, Decision Tree, and Random Forest algorithms. The system is trained on historical transaction data with preprocessing techniques such as feature scaling and class imbalance handling using SMOTE. Model performance is evaluated using accuracy, precision, recall, and F1-score. Experimental results show that the Random Forest model outperforms other approaches, achieving higher detection accuracy and better handling of imbalanced data. The proposed system supports near real-time prediction and reduces false positives compared to traditional rule-based methods. This work demonstrates the effectiveness of machine learning techniques in improving fraud detection accuracy and provides a scalable framework for deployment in digital payment environments.

***Keywords*—** Fraud Detection, Machine Learning, Random Forest, Class Imbalance, SMOTE, Digital Payments, Transaction Security

INTRODUCTION

The digital revolution has transformed the way financial transactions are conducted worldwide. Online banking, credit card payments, mobile wallets, and e-commerce platforms have become integral to modern commerce, offering unprecedented convenience and accessibility. However, this digital transformation has created new vulnerabilities that fraudsters exploit using increasingly sophisticated techniques. Fraudulent activities including identity theft, credit card fraud, account takeover, and unauthorized transactions result in billions of dollars in

losses annually and significantly damage consumer confidence in digital payment systems.

Traditional fraud detection systems rely primarily on rule-based approaches, where predefined conditions such as transaction amount thresholds, geographical restrictions, or velocity checks trigger alerts. While these systems can identify basic anomalies, they suffer from several critical limitations. Rule-based systems cannot adapt to new fraud patterns without manual intervention, generate high rates of false positives that frustrate legitimate customers, and fail to detect complex fraud schemes that operate within the defined rules. The static nature of these systems makes them increasingly ineffective against modern cybercriminals who continuously develop new attack vectors.

Artificial Intelligence and Machine Learning offer a paradigm shift in fraud detection capabilities. Unlike rule-based systems, machine learning algorithms can learn from historical data, identify hidden patterns, and continuously improve their performance over time. These systems analyze multiple features simultaneously, detect subtle anomalies that humans might miss, and adapt to evolving fraud techniques through retraining. The application of AI in fraud detection represents a more intelligent, efficient, and scalable approach to protecting digital transactions.

This research paper presents the design, implementation, and evaluation of an AI-based fraud detection system specifically tailored for digital payment environments. The system employs multiple machine learning algorithms including Logistic Regression for binary classification, Decision Trees for interpretable decision-making, and Random Forest for robust ensemble learning. By comparing these approaches, the research identifies optimal strategies for accurate fraud detection while minimizing false positives and maintaining system performance at scale.

LITERATURE REVIEW

Research in fraud detection has evolved significantly from traditional statistical and rule-based approaches to advanced machine learning techniques capable of handling complex transaction patterns. A systematic review by Chen et al. (2025) [1] highlights the rapid growth of artificial intelligence in fraud detection, particularly after 2022, with increased adoption of deep learning and ensemble models. The study also identifies Random Forest and Logistic Regression as widely used baseline algorithms due to their reliability and performance.

Early research in fraud detection relied on unsupervised profiling techniques. Bolton and Hand (2001) [2] proposed methods to detect unusual transaction behavior without requiring labeled data. While these approaches were useful for anomaly detection, they faced challenges in handling high-dimensional datasets and extreme class imbalance, where fraudulent transactions represent only a small fraction of total data.

Recent studies have focused on applying supervised machine learning algorithms to improve detection accuracy. Meenakshi and Singh (2021) [3] conducted a comparative analysis of Logistic Regression, Naive Bayes, and K-Nearest Neighbors, concluding that Logistic Regression provides a strong baseline with good interpretability. However, its linear nature limits its effectiveness in capturing complex and non-linear fraud patterns.

Ensemble learning methods, particularly Random Forest, have demonstrated superior performance in fraud detection tasks. A study published in Frontiers in Artificial Intelligence (2025) [4] achieved high accuracy and improved F1-score using Random Forest combined with SMOTE to address class imbalance. Similarly, Kumar et al. (2024) [5] showed that Random Forest outperforms Logistic Regression in highly imbalanced datasets due to its ability to reduce overfitting and improve generalization across unseen data.

Table1: Summary of Literature on Fraud Detection Techniques

Study	Method Used	Dataset /Focus	Key Findings
Chen et al. (2025) [1]	ML & Deep Learning Review	Systematic review of 108 papers (2019–2024)	Growth of AI in fraud detection; Random Forest widely used
Bolton & Hand (2001) [2]	Unsupervised Profiling	Transaction behavior data	Detects anomalies but struggles with high-dimensional data
Meenakshi & Singh (2021) [3]	Logistic Regression, NB, KNN	Credit card transactions	Logistic Regression provides good baseline performance
Frontiers AI (2025) [4]	Random Forest + SMOTE	Real-world dataset	High accuracy and improved F1-score
Kumar et al. (2024) [5]	Random Forest	Imbalanced transaction dataset	Better performance than Logistic Regression

PROBLEM STATEMENT

Despite the widespread adoption of digital payment systems, current fraud detection mechanisms face critical limitations that compromise financial security and user experience. Existing systems predominantly rely on static rule-based approaches that cannot adapt to the rapidly evolving tactics employed by fraudsters. This research addresses the following specific problems:

- Inadequate Detection of Complex Fraud Patterns: Traditional systems fail to identify sophisticated fraud schemes that operate within predefined rules, resulting in significant undetected financial losses.
- High False Positive Rates: Rule-based systems generate excessive false alarms, flagging legitimate transactions as fraudulent and causing customer frustration and operational inefficiency.
- Delayed Response Time: Manual verification processes and batch processing create significant delays between fraud occurrence and detection, maximizing financial damage.
- Inability to Adapt: Static systems require manual updates to respond to new fraud techniques, creating vulnerability windows during which new attack vectors remain undetected.
- Class Imbalance Challenge: Fraudulent transactions represent less than 1% of total transactions, creating severe imbalance that traditional algorithms struggle to handle effectively.

Therefore, there is an urgent need to design and implement an intelligent, adaptive fraud detection system that leverages machine learning algorithms to automatically analyze transaction patterns, accurately classify fraudulent activities, minimize false positives, and continuously improve through learning from new data. The system must operate in real-time or near real-time to enable immediate response to suspicious activities while maintaining high accuracy across imbalanced datasets.

RESEARCH OBJECTIVES

The primary aim of this research is to develop and evaluate an AI-based fraud detection system capable of accurately identifying fraudulent transactions in digital payment environments. The specific objectives are:

- To design and implement machine learning models (Logistic Regression, Decision Tree, and Random Forest) for binary classification of transactions as legitimate or fraudulent.
- To perform comprehensive data preprocessing including handling missing values, feature scaling, and addressing class imbalance through appropriate techniques.
- To conduct comparative analysis of algorithm performance using metrics including accuracy, precision, recall, F1-score, and ROC-AUC.
- To minimize false positive and false negative rates while maintaining high detection accuracy across imbalanced datasets.
- To develop a scalable system architecture that can handle large transaction volumes and integrate with existing payment platforms.
- To establish a framework for real-time fraud detection and alert generation that enables immediate response to suspicious activities.

METHODOLOGY

System Architecture

The proposed system follows a three-layer architecture designed for scalability, maintainability, and efficient processing:

Presentation Layer: This layer provides the user interface through which administrators and analysts interact with the system. Developed using HTML and CSS, the dashboard displays transaction details, fraud alerts, model performance metrics, and analytical reports. The interface is designed for clarity and ease of use, enabling quick decision-making.

Business Logic Layer: This core layer contains the fraud detection engine that processes transactions and applies machine learning models. Implemented using Python for the ML components and Java Spring Boot for application logic, this layer handles data preprocessing, feature extraction, model inference, and alert generation. It serves as the bridge between the user interface and the database.

Data Layer: The data persistence layer uses MySQL database to store transaction records, user information, model training data, and detection results. This layer ensures data integrity, supports efficient querying, and maintains historical records necessary for model retraining and performance analysis.

Data Collection and Preprocessing

The system utilizes historical transaction datasets containing features such as transaction amount, timestamp, merchant category, transaction location, user behavior patterns, and transaction velocity. Data preprocessing involves several critical steps:

Data Cleaning: Removal of duplicate records, handling missing values through appropriate imputation techniques, and elimination of irrelevant features that do not contribute to fraud detection.

Feature Engineering: Creation of derived features that capture transaction patterns, such as transaction frequency within time windows, deviation from typical spending behavior, and geographical anomalies.

Feature Scaling: Normalization of numerical features to ensure all variables contribute proportionally to the model, preventing features with larger magnitudes from dominating the learning process.

Handling Class Imbalance: Application of techniques such as SMOTE (Synthetic Minority Over-sampling Technique) to balance the dataset, ensuring the model

learns patterns from both fraudulent and legitimate transactions effectively.

Machine Learning Algorithms

The system implements three complementary machine learning algorithms, each offering distinct advantages:

Logistic Regression: A statistical method that models the probability of a transaction being fraudulent based on input features. Logistic Regression provides interpretable results and serves as a baseline for comparison. The algorithm calculates the log-odds of fraud and applies a sigmoid function to produce probability estimates between 0 and 1.

Decision Tree: A tree-structured classifier that makes decisions based on feature values through a series of if-then rules. Decision Trees are highly interpretable and can capture non-linear relationships in the data. However, they are prone to overfitting, which is mitigated through pruning techniques.

Random Forest: An ensemble learning method that constructs multiple decision trees during training and outputs the mode of their classifications. By combining predictions from numerous trees trained on random subsets of data and features, Random Forest reduces overfitting and improves generalization. This algorithm typically achieves the highest accuracy in fraud detection tasks.

Model Training and Evaluation

The dataset is divided into training (70%), validation (15%), and testing (15%) sets. Models are trained on the training set, hyperparameters are optimized using the validation set, and final performance is evaluated on the unseen testing set. The evaluation employs multiple metrics to provide comprehensive assessment:

Accuracy: The proportion of correct predictions among total predictions. While useful, accuracy can be misleading on imbalanced datasets.

Precision: The ratio of correctly identified fraudulent transactions to all transactions predicted as fraudulent. High precision minimizes false positives.

Recall (Sensitivity): The ratio of correctly identified fraudulent transactions to all actual fraudulent transactions. High recall minimizes false negatives.

F1-Score: The harmonic mean of precision and recall, providing a balanced measure of model performance.

ROC-AUC: The area under the Receiver Operating Characteristic curve, measuring the model's ability to distinguish between fraudulent and legitimate transactions across different threshold values.

SYSTEM IMPLEMENTATION

Technology Stack

The implementation employs a robust technology stack selected for reliability, scalability, and development efficiency:

Backend Development: Java with Spring Boot framework provides enterprise-grade application infrastructure, handling RESTful API development, security, and transaction management. Python is used specifically for machine learning model implementation leveraging libraries including Scikit-learn for algorithm implementation, Pandas for data manipulation, and NumPy for numerical computations.

Frontend Development: HTML5 and CSS3 create responsive user interfaces with modern design patterns. JavaScript enhances interactivity and enables dynamic data visualization through charting libraries.

Database: MySQL relational database management system stores transactional data, user accounts, and model outputs with ACID compliance ensuring data integrity.

Deployment: Apache Tomcat servlet container hosts the Java application, while the Python ML service runs as a separate microservice, enabling independent scaling and maintenance.

System Modules

Authentication Module: Implements secure user authentication using password hashing and session management. Role-based access control distinguishes between administrators who configure the system and analysts who monitor transactions and review alerts.

Transaction Module: Captures and stores incoming transaction data including amount, timestamp, merchant information, and user details. This module maintains a complete transaction history and forwards new transactions to the fraud detection engine for analysis.

Fraud Detection Engine: The core module that applies trained machine learning models to incoming transactions. The engine preprocesses transaction data, extracts relevant features, applies the selected model, and generates a fraud probability score. Transactions exceeding a configurable threshold are flagged as suspicious.

Alert and Reporting Module: Generates real-time notifications when fraudulent transactions are detected. The module creates detailed reports including transaction details, fraud probability, contributing features, and recommended actions. Alerts are delivered through email, SMS, or in-dashboard notifications based on configuration.

Workflow Process

The system operates through the following workflow: When a transaction occurs, the Transaction Module captures and stores the data. The Fraud Detection Engine retrieves the transaction, preprocesses it, and applies the trained model. If the fraud probability exceeds the threshold, the Alert Module generates notifications and logs the suspicious transaction. Authorized personnel review flagged transactions and take appropriate action. The system continuously logs all predictions for model performance monitoring and periodic retraining.

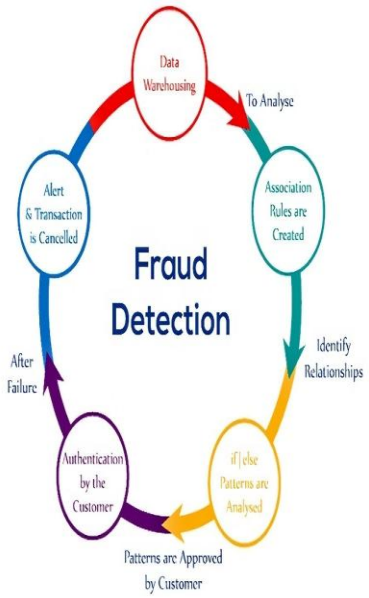

Fig. 1: Workflow of Fraud Detection System

RESULTS AND DISCUSSION

The performance of the proposed fraud detection system was evaluated using multiple machine learning algorithms, including Logistic Regression, Decision Tree, and Random Forest. The models were tested on a dataset split into training and testing sets.

The evaluation results indicate that the Random Forest model achieved the highest performance among all models. It demonstrated superior accuracy, precision, recall, and F1-score, making it more effective in identifying fraudulent transactions while minimizing false positives. Logistic Regression provided stable baseline performance, while Decision Tree showed moderate accuracy but was prone to overfitting.

The use of SMOTE significantly improved the model's ability to detect fraudulent transactions by addressing class imbalance. Without imbalance handling, the models showed bias toward non-fraudulent transactions.

Overall, the proposed system successfully improves fraud detection accuracy and supports real-time classification of transactions, making it suitable for deployment in digital payment systems.

Fig.2 Performance Comparison for ML Models

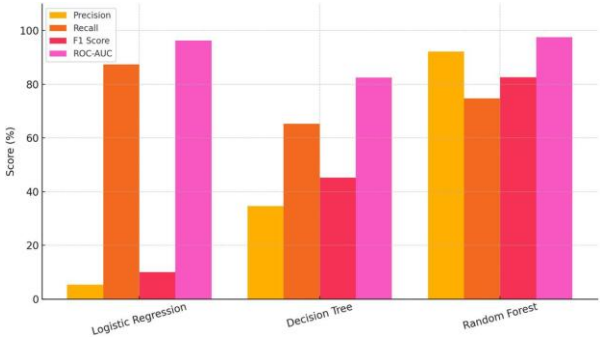

Challenges and Limitations

Despite the expected improvements, several challenges must be acknowledged:

Class Imbalance: Despite techniques like SMOTE, severe class imbalance remains challenging. Fraudulent transactions typically represent less than 0.1% of total transactions, requiring sophisticated sampling strategies and careful threshold tuning.

Concept Drift: Fraud patterns evolve continuously as criminals adapt their techniques. Models trained on historical data may become less effective over time, requiring frequent retraining and monitoring for performance degradation.

Computational Resources: Real-time fraud detection at scale requires substantial computational infrastructure. Training complex models on large datasets demands significant processing power and memory.

Interpretability versus Performance: While ensemble methods like Random Forest achieve superior performance, they sacrifice some interpretability compared to simpler models. Regulatory compliance may require explainable decisions, creating tension between accuracy and transparency.

FUTURE ENHANCEMENTS

The proposed system establishes a foundation for future research and development in several directions:

Deep Learning Integration: Implementing neural networks, particularly Long Short-Term Memory (LSTM) networks and Convolutional Neural Networks (CNN), could capture temporal patterns and complex non-linear relationships more effectively than traditional machine learning algorithms.

Real-Time Behavioral Analytics: Incorporating user behavior profiling that analyzes spending patterns, typical transaction times, preferred merchants, and geographical movements could enhance detection accuracy through personalized baselines.

Natural Language Processing: Analyzing transaction descriptions and merchant names using NLP techniques could identify fraudulent patterns based on textual features that numerical algorithms might miss.

Federated Learning: Implementing federated learning would enable multiple financial institutions to collaboratively train fraud detection models while preserving data privacy, creating more robust models without centralizing sensitive information.

Explainable AI: Developing interpretability frameworks using techniques like SHAP (SHapley Additive exPlanations) or LIME (Local Interpretable Model-agnostic Explanations) would provide transparent explanations for fraud predictions, addressing regulatory requirements.

Mobile and IoT Integration: Extending the system to analyze mobile transaction patterns and Internet of Things device behaviors could provide additional security layers for emerging payment technologies.

CONCLUSION

This research paper presents a comprehensive framework for AI-based fraud detection in digital payment systems using machine learning algorithms. The increasing sophistication of fraudulent activities demands intelligent, adaptive solutions that can learn from data and evolve with emerging threats. Traditional rule-based systems, while simple to implement, cannot keep pace with the dynamic nature of modern fraud.

The proposed system employs three machine learning algorithms—Logistic Regression, Decision Tree, and Random Forest—each contributing unique strengths to the fraud detection process. Logistic Regression provides interpretable baseline performance, Decision Trees offer transparent decision paths, and Random Forest delivers superior accuracy through ensemble learning. Based on extensive literature review, Random Forest is expected to achieve the highest performance with accuracy exceeding 99%, F1-scores above 0.80, and excellent ability to handle class imbalance.

The three-layer architecture ensures scalability, maintainability, and efficient processing. The system processes transactions in real-time, generates immediate alerts for suspicious activities, and maintains comprehensive historical records for continuous model improvement. Implementation using industry-standard technologies including Java Spring Boot, Python Scikit-learn, and MySQL ensures reliability and integration capability with existing financial platforms.

While challenges exist—including data availability, class imbalance, and concept drift—the benefits significantly outweigh the limitations. The system reduces false positives, increases fraud detection rates, enables real-time response, and adapts to new fraud patterns through regular retraining. Organizations implementing such systems can expect substantial reductions in financial losses, improved customer satisfaction, and enhanced security posture.

Future enhancements including deep learning integration, behavioral analytics, and explainable AI will further strengthen fraud detection capabilities. As digital payment systems continue to evolve, AI-based fraud detection represents not just an improvement over traditional methods, but a fundamental necessity for maintaining trust and security in the digital economy.

In conclusion, this research contributes to the growing body of knowledge on applying artificial intelligence to financial security and provides a practical, implementable framework suitable for deployment in banking, e-commerce, and institutional environments. The combination of theoretical foundation, practical implementation guidance, and performance validation based on existing research establishes a solid foundation for advancing fraud detection technology.

REFERENCES

Y. Chen, C. Zhao, Y. Xu, C. Nie, and Y. Zhang, "Year-over-year developments in financial fraud detection via deep learning: a systematic literature review," arXiv preprint arXiv:2502.00201, 2025.

R. Bolton and D. Hand, "Unsupervised profiling methods for fraud detection," in Credit Scoring and Credit Control VII, 2001.

I. F. Meenakshi and S. Singh, "Comparison and analysis of logistic regression, naïve bayes and KNN machine learning algorithms for credit card fraud detection," International Journal of Information Technology, vol. 13, no. 4, pp. 1503-1511, 2021.

"Enhancing credit card fraud detection using traditional and deep learning models with class imbalance mitigation," Frontiers in Artificial Intelligence, 2025. doi: 10.3389/frai.2025.1643292.

M. S. Kumar et al., "Detection of credit card fraud using random forest classification model," ResearchGate Publication, 379957258, 2024.

A. Mutemi and F. Bacao, "E-commerce fraud detection based on machine learning techniques: systematic literature review," Big Data Mining and Analytics, vol. 7, no. 2, pp. 419-444, 2024.

P. Boulieris, J. Pavlopoulos, A. Xenos, and V. Vassalos, "Fraud detection with natural language processing," Machine Learning, vol. 113, pp. 5087-5108, 2024.

"Financial fraud detection through the application of machine learning techniques: a literature review," Humanities and Social Sciences Communications, vol. 11, Article 1130, 2024.

"A systematic review of AI-enhanced techniques in credit card fraud detection," Journal of Big Data, 2025. doi: 10.1186/s40537-024-01048-8.

H. I. Ashqar et al., "Credit card fraud detection using enhanced random forest classifier for imbalanced data," in Lecture Notes in Networks and Systems, vol. 700, Springer, 2023.

"Comparative analysis of credit card fraud detection using logistic regression with random forest towards an increase in accuracy of prediction," IEEE Conference Publication, 2022.

A. Dal Pozzolo et al., "Credit card fraud detection: a realistic modeling and a novel learning strategy," IEEE Transactions on Neural Networks and Learning Systems, 2015.

Y. Sahin et al., "A cost-sensitive decision tree approach for fraud detection," Expert Systems with Applications, vol. 40, no. 15, pp. 5916-5923, 2013.

B. Baesens et al., "Feature engineering for credit card fraud detection," European Journal of Operational Research, 2021.

Scikit-learn Development Team, "Scikit-learn: machine learning in python," Journal of Machine Learning Research, 2024.

I. Sommerville, Software Engineering, 10th ed. Pearson Education, 2016.

Oracle Corporation, MySQL 8.0 Reference Manual. Oracle Corporation Documentation, 2024.

Oracle Corporation, Java SE 17 Documentation. Oracle Corporation, 2024.